

```
In [92]: #importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn import model_selection
from sklearn.model_selection import train_test_split, KFold, cross_val_score, StratifiedKFold, learning_curve
from sklearn.metrics import confusion_matrix, make_scorer, accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score
from sklearn.pipeline import Pipeline
from keras import models
model = models.Sequential()

from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from keras.utils import to_categorical
from scikeras.wrappers import KerasClassifier, KerasRegressor
```

```
In [93]: #Loading images and their labels
X = np.load('X.npy') # images
Y = np.load('Y.npy') # labels for the images (0 = no IDC, 1 = IDC)
```

```
In [94]: #making sure the data for X crossed over correctly
print(X[:3])
```

```
[[[226 164 206]
 [224 154 196]
 [225 175 211]
 ...
 [240 221 237]
 [232 184 214]
 [243 213 235]]

 [[217 142 188]
 [221 130 179]
 [224 150 196]
 ...
 [227 170 204]
 [229 180 215]
 [236 212 232]]

 [[237 178 212]
 [229 157 199]
 [218 125 175]]
```

```
In [95]: #making sure the data for Y crossed over correctly
print(Y[:5])

[0 0 0 0 0]
```

```
In [96]: #total number of images
print('Total number of images: {}'.format(len(X)))

Total number of images: 5547
```

```
In [97]: #total number of negative IDC images
print('Number of negative IDC Images: {}'.format(np.sum(Y==0)))

Number of negative IDC Images: 2759
```

```
In [98]: #total number of positive IDC images
print('Number of positive IDC Images: {}'.format(np.sum(Y==1)))
```

Number of positive IDC Images: 2788

```
In [99]: #shape of the images
print('Image shape (Width, Height, Channels): {}'.format(X[0].shape))
```

Image shape (Width, Height, Channels): (50, 50, 3)

```
In [100]: #train/test split
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2)
```

```
In [101]: # Reduce Sample Size
x_train = x_train[0:30000]
y_train = y_train[0:30000]
x_test = x_test[0:30000]
y_test = y_test[0:30000]
```

```
In [102]: # rescale pizel intensity
x_train = x_train / 256.0
x_test = x_test / 256.0
```

```
In [103]: #verifying shape
print("Training Data Shape:", x_train.shape)
print("Testing Data Shape:", x_test.shape)
```

Training Data Shape: (4437, 50, 50, 3)

Testing Data Shape: (1110, 50, 50, 3)

```
In [104]: #displaying the first five images in the training set along with the labels
for i in range(5):
    plt.imshow(x_train[i], cmap='gray'), plt.axis("off")
    plt.title('IDC = %d'%y_train[i])
    plt.show()
```

IDC = 0



```
In [105]: #printing the data to show that it's now 0-1
print(x_train[:3])
```

```
[[[0.9453125  0.94140625 0.94921875]
 [0.953125   0.92578125 0.94140625]
 [0.953125   0.94140625 0.9375   ]
 ...
 [0.94921875 0.9375   0.9453125 ]
 [0.9453125  0.94921875 0.953125 ]
 [0.94921875 0.9296875  0.93359375]]

 [[0.95703125 0.92578125 0.93359375]
 [0.9375     0.94140625 0.93359375]
 [0.95703125 0.92578125 0.94140625]
 ...
 [0.94921875 0.9375   0.9453125 ]
 [0.9609375  0.93359375 0.9453125 ]
 [0.9453125  0.953125  0.94140625]]

 [[0.94921875 0.9375   0.9453125 ]
 [0.953125   0.9296875 0.95703125]
 [0.94921875 0.93359375 0.94140625]]
```

```
In [106]: # reshape data
```

```
x_train_r = x_train.reshape(x_train.shape[0], x_train.shape[1]*x_train.shape[2]*x_train.shape[3])
x_test_r = x_test.reshape(x_test.shape[0], x_test.shape[1]*x_test.shape[2]*x_test.shape[3])

print("x_train shape: ",x_train_r.shape)
print("x_test shape: ",x_test_r.shape)
```

```
x_train shape: (4437, 7500)
x_test shape: (1110, 7500)
```

```
In [107]: def buildclassifier():
    classifier = Sequential() # initialize neural network
    classifier.add(Dense(units = 8, kernel_initializer = 'uniform', activation = 'relu', input_dim = x_train_r.shape[1]))
    classifier.add(Dense(units = 8, kernel_initializer = 'uniform', activation = 'relu'))
    classifier.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))
    classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
    return classifier
```

```
In [108]: classifier = KerasClassifier(build_fn = buildclassifier, epochs = 200)
accuracies = cross_val_score(estimator = classifier, X = x_train_r, y = y_train, cv = 6)
mean = accuracies.mean()
variance = accuracies.std()
```

```
116/116 [=====] - 1s 6ms/step - loss: 0.5137 - accuracy: 0.7591
Epoch 15/200
116/116 [=====] - 1s 6ms/step - loss: 0.5236 - accuracy: 0.7491
Epoch 16/200
116/116 [=====] - 1s 7ms/step - loss: 0.5227 - accuracy: 0.7493
Epoch 17/200
116/116 [=====] - 1s 6ms/step - loss: 0.5052 - accuracy: 0.7691
Epoch 18/200
116/116 [=====] - 1s 6ms/step - loss: 0.5080 - accuracy: 0.7647
Epoch 19/200
116/116 [=====] - 1s 7ms/step - loss: 0.5178 - accuracy: 0.7572
Epoch 20/200
116/116 [=====] - 1s 7ms/step - loss: 0.5187 - accuracy: 0.7555
Epoch 21/200
116/116 [=====] - 1s 6ms/step - loss: 0.5014 - accuracy: 0.7699
Epoch 22/200
116/116 [=====] - 1s 6ms/step - loss: 0.5076 - accuracy: 0.7610
Epoch 23/200
116/116 [=====] - 1s 6ms/step - loss: 0.5124 - accuracy: 0.7593
Epoch 24/200
```

```
In [109]: print("Accuracy mean: "+ str(mean))
          print("Accuracy variance: "+ str(variance))
```

Accuracy mean: 0.7189570152019407  
Accuracy variance: 0.027899708347895104

```
In [110]: #classifying decision tree and fitting it
          dtc = DecisionTreeClassifier()
          dtc_y = dtc.fit(x_train_r,y_train)
```

```
In [111]: dtc_pred = dtc_y.predict(x_test_r)
          print('Accuracy Score : ' + str(accuracy_score(y_test,dtc_pred)))
          print('Precision Score : ' + str(precision_score(y_test,dtc_pred)))
          print('Recall Score : ' + str(recall_score(y_test,dtc_pred)))
          print('F1 Score : ' + str(f1_score(y_test,dtc_pred)))
```

Accuracy Score : 0.6855855855855856  
Precision Score : 0.7064220183486238  
Recall Score : 0.6707317073170732  
F1 Score : 0.6881143878462913

```
In [112]: print('Confusion Matrix : \n' + str(confusion_matrix(y_test,dtc_pred)))
```

Confusion Matrix :  
[[376 160]  
 [189 385]]

In [ ]:

```
In [113]: #classifying random forest tree and fitting it
          rfc= RandomForestClassifier(n_estimators = 100, random_state=42)
          rfc_y = rfc.fit(x_train_r,y_train)
```

```
In [114]: rfc_pred = rfc_y.predict(x_test_r)
          print('Accuracy Score : ' + str(accuracy_score(y_test,rfc_pred)))
          print('Precision Score : ' + str(precision_score(y_test,rfc_pred)))
          print('Recall Score : ' + str(recall_score(y_test,rfc_pred)))
          print('F1 Score : ' + str(f1_score(y_test,rfc_pred)))
```

Accuracy Score : 0.7747747747747747  
Precision Score : 0.7852112676056338  
Recall Score : 0.7770034843205574  
F1 Score : 0.7810858143607706

```
In [115]: print('Confusion Matrix : \n' + str(confusion_matrix(y_test,rfc_pred)))
```

Confusion Matrix :  
[[414 122]  
 [128 446]]

In [ ]:

```
In [116]: #SVC classifying and fitting
          svc = SVC(random_state=42)
          svc_y = svc.fit(x_train_r,y_train)
```

```
In [117]: svc_pred = svc_y.predict(x_test_r)
          print('Accuracy Score : ' + str(accuracy_score(y_test,svc_pred)))
          print('Precision Score : ' + str(precision_score(y_test,svc_pred)))
          print('Recall Score : ' + str(recall_score(y_test,svc_pred)))
          print('F1 Score : ' + str(f1_score(y_test,svc_pred)))
```

Accuracy Score : 0.7864864864864864  
Precision Score : 0.8025134649910234  
Recall Score : 0.7787456445993032  
F1 Score : 0.790450928381963

```
In [118]: print('Confusion Matrix : \n' + str(confusion_matrix(y_test,svc_pred)))
```

```
Confusion Matrix :
[[426 110]
 [127 447]]
```

```
In [ ]:
```

```
In [119]: #Logistic regression classifying and fitting
```

```
lr = LogisticRegression()
lr_y = lr.fit(x_train_r,y_train)
```

C:\Users\breaan\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\linear\_model\\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

```
In [120]: lr_pred = lr_y.predict(x_test_r)
print('Accuracy Score : ' + str(accuracy_score(y_test,lr_pred)))
print('Precision Score : ' + str(precision_score(y_test,lr_pred)))
print('Recall Score : ' + str(recall_score(y_test,lr_pred)))
print('F1 Score : ' + str(f1_score(y_test,lr_pred)))
```

```
Accuracy Score : 0.7018018018018019
Precision Score : 0.7237569060773481
Recall Score : 0.6846689895470384
F1 Score : 0.7036705461056401
```

```
In [121]: print('Confusion Matrix : \n' + str(confusion_matrix(y_test,lr_pred)))
```

```
Confusion Matrix :
[[386 150]
 [181 393]]
```

```
In [ ]:
```

```
In [122]: #Logistic regression classifying and fitting
```

```
knn = KNeighborsClassifier()
knn_y = knn.fit(x_train_r,y_train)
```

```
In [123]: #Logistic regression accuracy
```

```
kscore = knn.score(x_test_r,y_test)
print("KNeighbors accuracy", kscore)
```

```
KNeighbors accuracy 0.7117117117117117
```

```
In [124]: knn_pred = knn_y.predict(x_test_r)
print('Accuracy Score : ' + str(accuracy_score(y_test,knn_pred)))
print('Precision Score : ' + str(precision_score(y_test,knn_pred)))
print('Recall Score : ' + str(recall_score(y_test,knn_pred)))
print('F1 Score : ' + str(f1_score(y_test,knn_pred)))
```

```
Accuracy Score : 0.7117117117117117
Precision Score : 0.7702127659574468
Recall Score : 0.6306620209059234
F1 Score : 0.6934865900383141
```

```
In [125]: print('Confusion Matrix : \n' + str(confusion_matrix(y_test,lr_pred)))
```

```
Confusion Matrix :
[[386 150]
 [181 393]]
```

```
In [ ]:
```

```
In [126]: results = []
results.append(mean)
results.append(dscore)
results.append(rscore)
results.append(sscore)
results.append(lscore)
results.append(kscore)
print(results)

[0.7189570152019407, 0.6891891891891891, 0.7738738738738739, 0.7801801801801802, 0.69009009009009, 0.7117117117117117]
```

```
In [127]: models = []
models.append(('classifier', KerasClassifier()))
models.append(('dtc', DecisionTreeClassifier()))
models.append(('rfc', RandomForestClassifier()))
models.append(('svc', SVC()))
models.append(('ls', LogisticRegression()))
models.append(('knn', KNeighborsClassifier()))
print(models)

[('classifier', KerasClassifier(
    model=None
    build_fn=None
    warm_start=False
    random_state=None
    optimizer=rmsprop
    loss=None
    metrics=None
    batch_size=None
    validation_batch_size=None
    verbose=1
    callbacks=None
    validation_split=0.0
    shuffle=True
    run_eagerly=False
    epochs=1
    class_weight=None
)), ('dtc', DecisionTreeClassifier()), ('rfc', RandomForestClassifier()), ('svc', SVC()), ('ls', LogisticRegression()), ('knn', KNeighborsClassifier())]
```

```
In [ ]:
```

```
In [ ]:
```

The below code is something I was working on; the gridsearch. But, couldn't get it to work for this project.

```
minmax = MinMaxScaler()
```

```
dtc = DecisionTreeClassifier()
rfc= RandomForestClassifier(n_estimators = 100, random_state=42)
svc = SVC(random_state=42)
lr = LogisticRegression()
knn = KNeighborsClassifier()
dtr = DecisionTreeRegressor()
```

```
pipe = Pipeline(steps = [('scaler', minmax), ('classifier', dtr)])
```

```
pipe.fit(x_train_r, y_train)
```

```
param_grid = [
    {'classifier_max_depth': [2,6,8,10],
     'classifier_min_samples_split': [2,5,10,15],
     'classifier': [dtr],
     'classifier_max_depth': [2,6,8,10],
     'classifier_min_samples_split': [2,5,10,15],
     'classifier_max_leaf_nodes': [None,10,20,50,100]},
    {'classifier': [rfc],
     'classifier_max_depth': [2,6,8,10],
     'classifier_min_samples_split': [2,5,10,15],
     'classifier_max_features': [2,3,4,5,6]},
    {'classifier': [svc],
     'classifier_max_depth': [2,6,8,10],
     'classifier_min_samples_split': [2,5,10,15],
     'classifier_max_features': [2,3,4,5,6]},
    {'classifier': [lr],
     'classifier_max_depth': [2,6,8,10],
     'classifier_min_samples_split': [2,5,10,15],
     'classifier_max_features': [2,3,4,5,6]},
    {'classifier': [knn],
     'classifier_max_depth': [2,6,8,10],
     'classifier_min_samples_split': [2,5,10,15],
     'classifier_max_features': [2,3,4,5,6]},
    {'classifier': [dtr],
     'classifier_max_depth': [2,6,8,10],
     'classifier_min_samples_split': [2,5,10,15],
     'classifier_max_features': [2,3,4,5,6]}
]
```

```
grid_search = GridSearchCV(pipe, param_grid, cv = 5, verbose = 0)
```

```
best_model = grid_search.fit(x_train_r, y_train)
```